

VMA Translation in Multi-process OS

LZMSCI521M | Week 1

Dr Nono Saha

What Will You Learn Today?

1. **Virtual Memory**
2. **Address Translation**
3. **Translation Lookaside Buffer (TLB)**
4. **Takeaway**
5. **Practical Exercises**
6. **What Next?**


1.1

Virtual Memory

For Second-year Computer Science Students

University of Leipzig/ ScaDS.AI

Why Virtual Memory?




Possible problems

- 1 Computers have limited physical memory, but modern programs expect a large address space.
- 2 Different programs (processes) need to be isolated: can't let one process overwrite another's memory!
- 3 The ISA often supports a larger address space than there is physical memory.

Virtual Memory Concepts

- ▶ **Idea:** Give the programmer the illusion of a large addressspace while having a small physical memory.
 - ▶ So that the programmer does not worry about managing physical memory.
- ▶ Programmer can assume having an “infinite” amount of physical memory
- ▶ Hardware and software cooperatively and automatically manage the physical memory space to provide the illusion
 - ▶ Illusion is maintained for each independent process

Example: Suppose a process wants to read byte 0x2400 (virtual).




Thus, virtual address 0x2400 maps to physical address 0x5400.

Example of Translation in Multi-Process Situation

Process A: Virtual 0x4000 (Page 4) → Frame 12 → 0xC000


Process B: Virtual 0x4000 (Page 4) → Frame 7 → 0x7000



Both processes use the same virtual address but access different physical locations.

- ▶ Indirection (in addressing)
- ▶ Address generated by each instruction in a program is a “virtual address”
 - ▶ i.e., it is not the physical address used to address main memory
 - ▶ called “linear address” in x86
- ▶ An “address translation” mechanism maps this address to a “physical address”
 - ▶ called “real address” in x86q
 - ▶ Address translation mechanism can be implemented in hardware and software together

A system with Virtual Memory (Page based)



- **Address Translation:** The hardware converts virtual addresses into physical addresses via an OS-managed lookup table (page table).

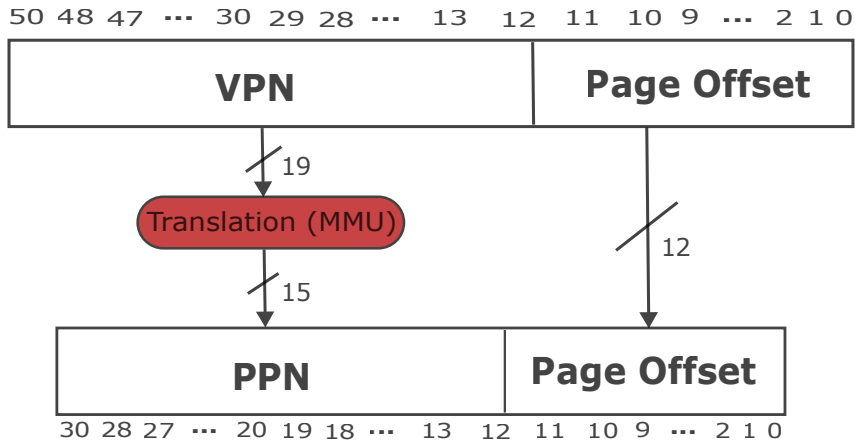
1.2

Address Translation

For Second-year Computer Science Students

University of Leipzig/ ScaDS.AI

Virtual Addresses



Physical Addresses

Virtual Memory Example

▶ System:

- ▶ Virtual Memory capacity: $2 \text{ GB} = 2^{31} \text{ bytes}$
- ▶ Physical Memory capacity: $128 \text{ MB} = 2^{27} \text{ bytes}$
- ▶ Page size: $4 \text{ KB} = 2^{12} \text{ bytes}$

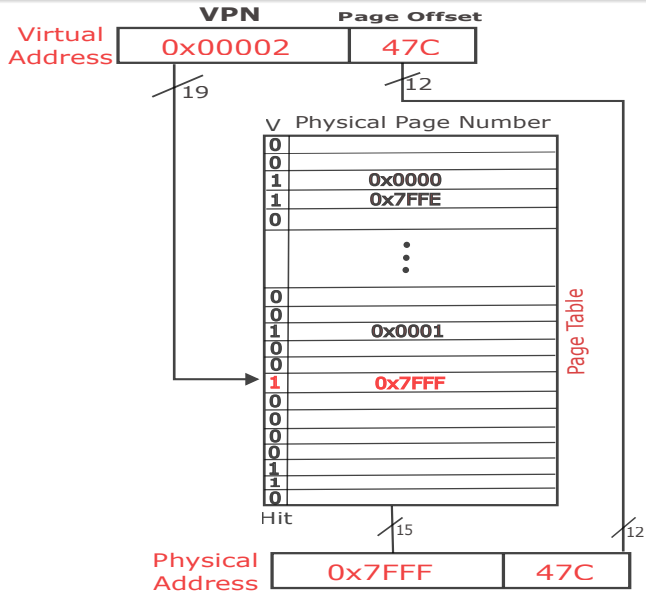
▶ Organisation:

- ▶ Virtual address: 31 bits
- ▶ Physical address: 27 bits
- ▶ Page offset: 12 bits
- ▶ Virtual pages $= 2^{31}/2^{12} = 2^{19}$ (VPN = 19 bits)
- ▶ Physical pages $= 2^{27}/2^{12} = 2^{15}$ (PPN = 15 bits)

How Do We translate Addresses?

- ▶ Page table
 - ▶ Has entry for each virtual page
- ▶ Each page table entry has:
 - ▶ Valid bit: whether the virtual page is located in physical memory (if not, it must be fetched from the hard disk)
 - ▶ Physical page number: where the virtual page is located in physical memory
 - ▶ (Replacement policy, dirty bits)

Page Table Example



Page Table Example (Homework 1)

Consider the PPN on the right.

- 1 What is the physical address of virtual address **0x5F20**?
- 2 What is the physical address of virtual address **0x73E0**?

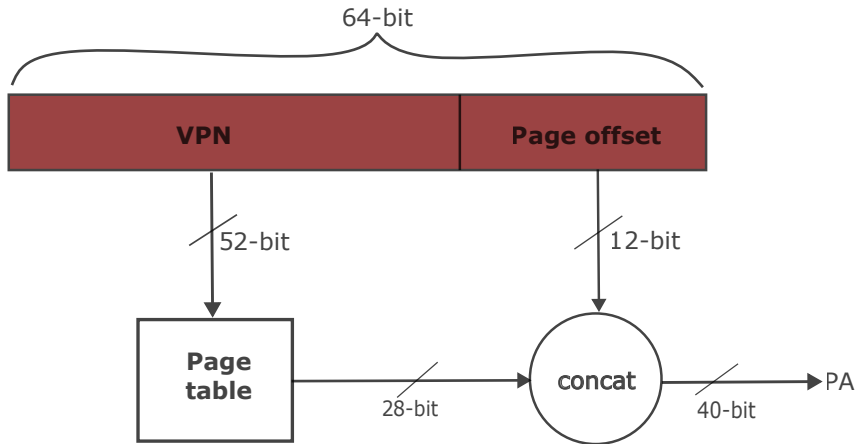
v Physical Page Number	
0	
0	
1	0x0000
1	0x7FFE
0	
	⋮
0	
0	
1	0x0001
0	
0	
1	0x7FFF
0	
0	
0	
0	
1	
1	
0	

Page Table

Hit

15

Issues with the Page Table Size



- Suppose a 64-bit VA and a 40-bit PA, how large is the page table?

Page Table Challenges

- ▶ Page table accesses have a lot of temporal locality
 - ▶ at least part of it needs to be located in physical memory
 - ▶ Data accesses have temporal and spatial locality
 - ▶ Large page size (say 4KB, 8KB, or even 1-2GB), so consecutive loads/stores likely to access same page
- ▶ Each load/store requires at least two memory accesses:
 - 1 one for address translation (page table read)
 - 2 one to access data with the physical address (after translation)
- ▶ Two memory accesses to service a load/store greatly degrades load/store execution time
 - ▶ Unless we are clever... 31

1.3

Translation Lookaside Buffer (TLB)

For Second-year Computer Science Students

University of Leipzig/ ScaDS.AI

Translation Lookaside Buffer (TLB)

Idea: Cache the page table entries (PTEs) in a hardware structure in the processor

Advantages:

- ▶ Small cache of most recently used translations (PTEs)
- ▶ Reduces number of memory accesses required for most loads/stores to only one
- ▶ High associativity
- ▶ Typically 16 - 512 entries
- ▶ > 95-99 % hit rates typical (depends on workload)

Memory Protection

- ▶ Multiple programs (processes) run at once
 - ▶ Each process has its own page table
 - ▶ Each process can use entire virtual address space without worrying about where other programs are
- ▶ A process can only access physical pages mapped in its page table—cannot overwrite memory of another process
 - ▶ Provides protection and isolation between processes
 - ▶ Enables access control mechanisms per page

Page Table is per Process

- ▶ Each process has its own virtual address space
 - ▶ Full address space for each program
 - ▶ Simplifies memory allocation, sharing, linking and loading.

1.4

Takeaway

For Second-year Computer Science Students

University of Leipzig/ ScaDS.AI

Key Takeaways

- ▶ **Virtual memory** gives each process a large, private address space.
- ▶ The **OS and hardware** jointly translate virtual addresses to physical addresses.
- ▶ **Page tables** do the mapping; **TLB** makes it fast.
- ▶ Per-process page tables enforce **memory protection**.
- ▶ A subset of virtual pages are located in physical memory

1.5

Practical Exercises

For Second-year Computer Science Students

University of Leipzig/ ScaDS.AI

- 1 Consider a virtual address space of 64 pages of 1024 words each, mapped onto a physical memory of 32 frames.
 - a. How many bits are there in the logical address?
 - b. How many bits are there in the physical address?
- 2 Assuming a 1-KB page size, what are the page numbers and offsets for the following address references (provided as decimal numbers):
 - a. 3085
 - b. 42095
 - c. 2000001
- 3 Suppose process X and process Y both have a virtual page 2.
 - a. X 's page 2 maps to frame 13; Y 's page 2 to frame 21.
 - b. If Y reads from virtual address 0x2004 (4KB pages), what physical address in RAM is accessed?

1.6

What Next?

For Second-year Computer Science Students

University of Leipzig/ ScaDS.AI

What next?

Page Fault and Page replacement, Structure of Page tables.

- 1 Cache Memory (Cache of Level 1, 2, 3 and 4)
- 2 Swapping Memory
- 3 IA-32 and x86-64 Memory Architecture

Some Important Materials

- ▶ Operating System Concepts, Abraham Silberschatz, Peter Baer G. and Greg Gagne, Tenth Edition, 2018.
- ▶ Youtube Course on Virtual Memory by Prof. Onur Mutlu, ETH Zürich
- ▶ Code and Course Material